

3. Reasoning as Memory

Introduction

Item memory


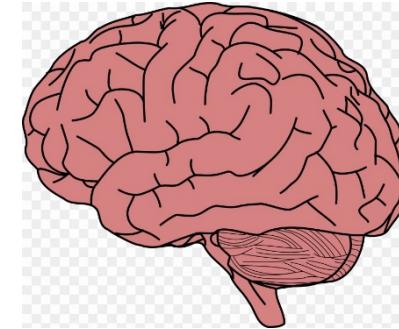
Relational memory

Program memory

Introduction


Memory is part of intelligence

- Memory is the ability to **store**, **retain** and **recall** information
- Brain memory stores items, events and high-level structures
- Computer memory stores data and temporary variables



Memory-reasoning analogy

- 2 processes: **fast-slow**
 - Memory: **familiarity-recollection**
- Cognitive test:
 - Corresponding reasoning and memorization performance
 - Increasing # premises, inductive/deductive reasoning is affected



Common memory activities


- **Encode**: write information to the memory, often requiring compression capability
- **Retain**: keep the information overtime. This is often assumed in machinery memory
- **Retrieve**: read information from the memory to solve the task at hand

Encode

Retain

Retrieve

Memory taxonomy based on memory content



Item memory

Associative memory

RAM-like memory

Independent memory

Distributed item memory as associative memory

Language

"Green" means
"go," but what
does "red" mean?

Time

birthday party on
30th Jan

Object

Where is my pen?
What is the
password?

Behaviour





Semantic
memory

Episodic
memory

Working
memory

Motor
memory

Associate memory can be implemented as Hopfield network



Encode $\hat{\mathbf{M}} = \sum_{k=1}^q \mathbf{b}_k \mathbf{a}_k^T$

Retrieve

$\mathbf{b} = \hat{\mathbf{M}} \mathbf{a}_j$

Feed-forward retrieval


Retrieve

$$x_i(t+1) = \text{sign} \left(\sum_{j=1}^N w_{ij} x_j(t) \right)$$


Recurrent retrieval

Rule-based reasoning with associative memory

- Encode a set of rules:
“pre-conditions
→post-conditions”
- Support variable binding, rule-conflict handling and partial rule input
- Example of encoding rule “A:1,B:3,C:4 → X”




Memory-augmented neural networks: computation-storage separation




Neural Turing Machine (NTM)

- Memory is a **2d matrix**
- Controller is a **neural network**
- The controller read/writes to memory at certain addresses.
- Trained **end-to-end**, differentiable
- Simulate Turing Machine
→ support symbolic reasoning, algorithm solving



Addressing mechanism in NTM




Optimal memory writing for memorization

- Simple finding: writing too often deteriorates memory content (not retainable)
- Given input sequence of length T and only D writes, *when should we write to the memory?*

Theorem 3. Given D memory slots, a sequence with length T , a decay rate $0 < \lambda \leq 1$, then the optimal intervals $\{l_i \in \mathbb{R}^+\}_{i=1}^{D+1}$ satisfying $T = \sum_{i=1}^{D+1} l_i$ such that the lower bound on the average contribution $I_\lambda = \frac{C}{T} \sum_{i=1}^{D+1} f_\lambda(l_i)$ is maximized are the following:

$$l_1 = l_2 = \dots = l_{D+1} = \frac{T}{D+1} \quad (7)$$




Uniform writing is optimal for memorization



Better memorization means better algorithmic reasoning

| | | |
|-----|---------------------|--|
| Max | $x_1 x_2 \dots x_T$ | $\max(x_1, x_2) \max(x_3, x_4) \dots \max(x_{T-1}, x_T)$ |
|-----|---------------------|--|


T=50, D=5



Regular



Uniform (cached)



Memory of independent entities

Weston, Jason, Bordes, Antoine, Chopra, Sumit, and Mikolov, Tomas.
Towards ai-complete question answering: A set of prerequisite toy tasks. CoRR, abs/1502.05698, 2015.

- Each slot store one or some entities
 - Memory writing is done separately for each memory slot
- each slot maintains the life of one or more entities
- The memory is a set of N parallel RNNs

Task 3: Three Supporting Facts

John picked up the apple.

John went to the office.


John went to the kitchen.

John dropped the apple.

Where was the apple before the kitchen? A:office



Recurrent entity network



Mary picked up the ball.

Mary went to the garden.

“Where is the ball?”

$$p_j = \text{Softmax}(q^T h_j)$$


$$u = \sum_j p_j h_j$$

$$y = R\phi(q + Hu)$$

Garden

Henaff, Mikael, Jason Weston, Arthur Szlam, Antoine Bordes, and Yann LeCun.
"Tracking the world state with recurrent entity networks."
In 5th International Conference on Learning Representations, ICLR 2017. 2017.


Recurrent Independent Mechanisms



$$\tilde{h}_{t,k} = D_k(h_{t,k}) = \text{LSTM}(h_{t,k}, A_k^{(in)}; \theta_k^{(D)}) \quad \forall k \in \mathcal{S}_t$$

$$Q_{t,k} = \tilde{W}_k^q \tilde{h}_{t,k}, \forall k \in \mathcal{S}_t \quad K_{t,k} = \tilde{W}_k^e \tilde{h}_{t,k}, \forall k \quad V_{t,k} = \tilde{W}_k^v \tilde{h}_{t,k}, \forall k$$

$$h_{t+1,k} = \text{softmax} \left(\frac{Q_{t,k}(K_{t,:})^T}{\sqrt{d_e}} \right) V_{t,:} + \tilde{h}_{t,k} \quad \forall k \in \mathcal{S}_t.$$




$$A_k^{(in)} = \text{softmax} \left(\frac{h_t W_k^q (X W^e)^T}{\sqrt{d_e}} \right) X W^v,$$

Relational memory

Graph memory


Tensor memory

Motivation for relational memory: item memory is weak at recognizing relationships




- Store and retrieve individual items
- Relate pair of items of the same time step
- Fail to **relate temporally distant items**

$$\hat{\mathbf{M}} = \sum_{k=1}^q \mathbf{b}_k \mathbf{a}_k^T$$



Dual process in memory




Howard Eichenbaum, *Memory, amnesia, and the hippocampal system* (MIT press, 1993).

Alex Konkel and Neal J Cohen, "Relational memory and the hippocampus: representations and methods", *Frontiers in neuroscience* 3 (2009).

Memory as graph

- Memory is a **static graph** with fixed nodes and edges
- Relationship is somehow known
- Each memory node stores the state of the graph's node
- Write to node via message passing
- Read from node via MLP




bAbl

$$x_i = \text{MLP}(\text{concat}(\text{last}(\text{LSTM}_S(s_i)), \text{last}(\text{LSTM}_Q(q)), \text{onehot}(p_i + o)))$$

Node

Fact 1



Fact 2

Fact 3

$$m_{ij}^t = f(h_i^{t-1}, h_j^{t-1}, e_{ij})$$

| Method | N | Mean Error (%) | Failed tasks (err. >5%) |
|------------------|----|-----------------|-------------------------|
| RRN* (this work) | 15 | 0.46 ± 0.77 | 0.13 ± 0.35 |

CLEVER

$$o_i = \text{concat}(p_i, \text{onehot}(c_i), \text{onehot}(m_i))$$


$$q = \text{concat}(\text{onehot}(s), \text{onehot}(n))$$

$$x_i = \text{MLP}(\text{concat}(o_i, q))$$

Node
(colour, shape, position)


Edge
(distance)

solution to the question: "green, 3 jumps", which is "plus",




Memory of graphs access conditioned on query

- Encode multiple graphs, each graph is stored in a set of memory row
- For each graph, the controller read/write to the memory:
 - Read uses content-based attention
 - Write use message passing
- Aggregate read vectors from all graphs to create output




Capturing relationship can be done via memory slot interactions using attention

- Graph memory needs customization to an explicit design of nodes and edges
- Can we automatically learn structure with a 2d tensor memory?
- Capture relationship: each slot interacts with all other slots (self-attention)



Relational Memory Core (RMC) operation



$$s_{i,t} = (h_{i,t-1}, m_{i,t-1})$$

$$f_{i,t} = W^f x_t + U^f h_{i,t-1} + b^f$$

$$i_{i,t} = W^i x_t + U^i h_{i,t-1} + b^i$$

$$o_{i,t} = W^o x_t + U^o h_{i,t-1} + b^o$$

$$m_{i,t} = \sigma(f_{i,t} + \tilde{b}^f) \circ m_{i,t-1} + \sigma(i_{i,t}) \circ \underbrace{g_\psi(\tilde{m}_{i,t})}_{\text{row/memory-wise MLP with layer normalisation}}$$


$$h_{i,t} = \sigma(o_{i,t}) \circ \tanh(m_{i,t})$$

$$s_{i,t+1} = (m_{i,t}, h_{i,t})$$


RNN-like Interface

$$\widetilde{M} = \text{softmax} \left(\frac{MW^q ([M; x]W^k)^T}{\sqrt{d^k}} \right) [M; x]W^v,$$



Allowing pair-wise interactions can answer questions on temporal relationship



What is the N^{th} farthest from vector m'




Dot product attention works for simple relationship, but ...




For hard relationship, scalar representation is limited

Self-attentive associative memory



Complicated relationship needs high-order relational memory




Program memory

Module memory

Stored-program memory


Predefining program for subtask

- A program designed for a task becomes a **module**
- Parse a question to module layout (order of program execution)
- Learn the weight of each module to master the task




Program selection is based on parser, others are end2end trained

5 module templates




The most powerful memory is one that stores both program and data

- Computer architecture: Universal Turing Machines/Harvard/VNM
- **Stored-program principle**
- Break a big task into subtasks, each can be handled by a TM/single purposed program stored in a program memory



NUTM: Learn to select program (neural weight) via program attention


- Neural stored-program memory (NSM) stores **key (the address)** and **values (the weight)**
- The weight is selected and loaded to the controller of NTM
- The stored NTM weights and the weight of the NUTM is learnt end-to-end by backpropagation



Le, Hung, Truyen Tran, and Svetha Venkatesh. "Neural Stored-program Memory." In *International Conference on Learning Representations*. 2019.

Scaling with memory of mini-programs

- Prior, 1 program = 1 neural network (millions of parameters)
- Parameter inefficiency since the programs do not share common parameters
- Solution: **store sharable mini-programs to compose infinite number of programs**



it is analogous to building Lego structures corresponding to inputs from basic Lego bricks.

Recurrent program attention to retrieve singular components of a program


$$\mathbf{P}_t = \mathbf{U} \mathbf{S} \mathbf{V}^T$$

$$= \sum_n^{r_m} \sigma_{tn} u_{tn} v_{tn}^\top$$

$$u_{tn} = \sum_{i=1}^{P_u} w_{tin}^u \mathbf{M}_U(i)$$



$$v_{tn} = \sum_{i=1}^{P_v} w_{tin}^v \mathbf{M}_V(i)$$

$$\sigma_{tn} = \begin{cases} \text{softplus} \left(\sum_{i=1}^{P_s} w_{tin}^\sigma \mathbf{M}_S(i) \right) & n = r_m \\ \sigma_{tn+1} + \text{softplus} \left(\sum_{i=1}^{P_s} w_{tin}^\sigma \mathbf{M}_S(i) \right) & n < r_m \end{cases}$$



Program attention is equivalent to
binary decision tree reasoning


Recurrent program attention auto
detects task boundary



(a)



(c)



(b)

(d)

QA

10. Combinatorics reasoning

RNN

MANN


GNN

Transformer

Implement combinatorial algorithms with neural networks

Processor P:

- (a) aligned with the computations of the target algorithm;
- (b) operates by matrix multiplications, hence natively admits useful gradients;
- (c) operates over high-dimensional latent spaces



Processor as RNN

- Do not assume knowing the structure of the input, input as a sequence

→ not really reasonable, harder to generalize

- RNN is Turing-complete → can simulate any algorithm


- But, it is not easy to learn the simulation from data (input-output)

→ Pointer network


Assume $O(N)$ memory
And $O(N^2)$ computation
 N is the size of input

$$u_j^i = v^T \tanh(W_1 e_j + W_2 d_i) \quad j \in (1, \dots, n)$$


$$p(C_i | C_1, \dots, C_{i-1}, \mathcal{P}) = \text{softmax}(u^i)$$




(a) Sequence-to-Sequence




(b) Ptr-Net



(d) Ptr-Net, $m=5-50$, $n=500$



(e) Ptr-Net, $m=50$, $n=50$




(f) Ptr-Net, $m=5-20$, $n=20$

Vinyals, Oriol, Meire Fortunato, and Navdeep Jaitly. "Pointer networks."



In *Proceedings of the 28th International Conference on Neural Information Processing Systems-Volume 2*, pp. 2692-2700. 2015.

Processor as MANN


- MANN simulates neural computers or Turing machine → ideal for implement algorithms
- Sequential input, no assumption on input structure
- Assume $O(1)$ memory and $O(N)$ computation



DNC: item memory for graph reasoning



NUTM: implementing multiple algorithms at once



Le, Hung, Truyen Tran, and Svetha Venkatesh. "Neural Stored-program Memory." In *International Conference on Learning Representations*. 2019.

STM: relational memory for graph reasoning





Figure 2. Learning curves on algorithmic synthetic tasks.

| Model | #Parameters | Convex hull | | TSP | | Shortest path | Minimum spanning tree |
|------------|--------------|--------------|--------------|---------------------|---------------------|---------------|-----------------------|
| | | $N = 5$ | $N = 10$ | $N = 5$ | $N = 10$ | | |
| LSTM | 4.5 M | 89.15 | 82.24 | 73.15 (2.06) | 62.13 (3.19) | 72.38 | 80.11 |
| ALSTM | 3.7 M | 89.92 | 85.22 | 71.79 (2.05) | 55.51 (3.21) | 76.70 | 73.40 |
| DNC | 1.9 M | 89.42 | 79.47 | 73.24 (2.05) | 61.53 (3.17) | 83.59 | 82.24 |
| RMC | 2.8 M | 93.72 | 81.23 | 72.83 (2.05) | 37.93 (3.79) | 66.71 | 74.98 |
| STM | 1.9 M | 96.85 | 91.88 | 73.96 (2.05) | 69.43 (3.03) | 93.43 | 94.77 |

Table 3. Prediction accuracy (%) for geometry and graph reasoning tasks with random *one-hot* associated features. Italic numbers are tour length–additional metric for TSP. Average optimal tour lengths found by brute-force search for $N = 5$ and 10 are 2.05 and 2.88, respectively.

Processor as graph neural network (GNN)




Motivation:

- Many algorithm operates on graphs
- Supervise graph neural networks with algorithm operation/step/final output
- Encoder-Process-Decode framework:

<https://petar-v.com/talks/Algo-WWW.pdf>


Veličković, Petar, Rex Ying, Matilde Padovano, Raia Hadsell, and Charles Blundell.

"Neural Execution of Graph Algorithms." In *International Conference on Learning Representations*. 2019.




Example: GNN for a specific problem (DNF counting)

- Count #assignments that satisfy disjunctive normal form (DNF) formula
- Classical algorithm is P-hard $O(mn)$
- m : #clauses, n : #variables
- Supervised training




$$\hat{v}_{x_c,t+1} = L_{c_1}\left(v_{x_c,t}, \sum_{x_l \in N(x_l)} M_l(v_{x_l,t})\right)$$




(a)

$$v_{x_d,t+1} = L_d\left(v_{x_d,t}, \sum_{x_c \in N(x_d)} M_c(\hat{v}_{x_c,t+1})\right)$$




(b)

$$v_{x_c,t+1} = L_{c_2}\left(\hat{v}_{x_c,t+1}, M_d(v_{x_d,t+1})\right).$$




(c)

$$v_{x_l,t+1} = L_l\left(v_{x_l,t}, \left(\sum_{x_c \in N(x_l)} M_c(v_{x_c,t+1}) \parallel M_l(v_{\neg x_l,t}) \right)\right)$$

Best: $O(m+n)$

Example: GNN trained with reinforcement learning (maximum common subgraph)

- Maximum common subgraph (MCS) is NP-hard
- Search for MCS:
 - BFS then pruning
 - Which node to visit first?
- Cast to RL:
 - State:
 - Current subgraph
 - Node-node mapping
 - Input graph
 - Action: Node pair or edge will be visited
 - Reward: +1 if a node pair is selected
 - $Q(s,a)$ =largest common subgraph size




Learning state representation with GNN


$$Q^*(s_t, a_t), \text{ as } r_t + \gamma V^*(s_{t+1})$$

$$Q(s_t, a_t) = 1 + \gamma \text{MLP} \left(\text{CONCAT} \left(\text{INTERACT}(\mathbf{h}_{\mathcal{G}_1}, \mathbf{h}_{\mathcal{G}_2}), \text{INTERACT}(\mathbf{h}_{s1}, \mathbf{h}_{s2}), \mathbf{h}_{\mathcal{D}c}, \mathbf{h}_{D_0} \right) \right).$$

Pretrain with ground-truth Q or expert estimation
Then train as DQN



(a) Result on ROAD-CA with 978513 nodes.




(b) Result on ROAD-TX with 1080909 nodes.

Bidomain representation

$$\begin{aligned} \mathbf{h}_{D_k} = \text{INTERACT} & \left(\text{READOUT}(\{\mathbf{h}_i | i \in \mathcal{V}'_{k1}\}), \right. \\ & \left. \text{READOUT}(\{\mathbf{h}_j | j \in \mathcal{V}'_{k2}\}) \right). \\ \text{READOUT}(\{\mathbf{h}_{D_k} | k \in \mathcal{D}^{(c)}\}). \end{aligned}$$


| Method | ROAD | DBEN |
|--|--------------|--------------|
| GLSEARCH (no $\mathbf{h}_{\mathcal{G}}$) | 0.977 | 0.878 |
| GLSEARCH (no \mathbf{h}_s) | 1.000 | 0.874 |
| GLSEARCH (no $\mathbf{h}_{\mathcal{D}c}$) | 0.803 | 0.780 |
| GLSEARCH (no \mathbf{h}_{D_0}) | 0.576 | 0.856 |
| GLSEARCH (SUM interact) | 0.902 | 0.913 |
| GLSEARCH (unfactored) | 0.447 | 0.807 |
| GLSEARCH (unfactored-i) | 0.500 | 0.789 |
| GLSEARCH | 0.992 | 1.000 |
| BEST SOLUTION SIZE | 132 | 508 |

Neural networks and algorithms alignment




Summary statistics

What is the maximum value difference among treasures?




Relational argmax

What are the colors of the furthest pair of objects?



Dynamic programming

What is the cost to defeat monster X by following the optimal path?



NP-hard problem

Subset sum: Is there a subset that sums to 0?



MLPs
~ feature extraction




Deep Sets (Zaheer et al., 2017)
~ summary statistics


$$y = \text{MLP}(\|_{s \in S} X_s)$$

$$y = \text{MLP}_2 \left(\sum_{s \in S} \text{MLP}_1(X_s) \right)$$

<https://petar-v.com/talks/Algo-WWW.pdf>



GNNs
~ (pairwise) relations



GNNs
~ (pairwise) relations

$$\begin{aligned} h_s^{(k)} &= \sum_{t \in S} \text{MLP}_1^{(k)} \left(h_s^{(k-1)}, h_t^{(k-1)} \right) \\ y &= \text{MLP}_2 \left(\sum_{s \in S} h_s^{(K)} \right) \end{aligned}$$

Neural exhaustive search

GNN is aligned with Dynamic Programming (DP)

Graph Neural Network

```
for k = 1 ... GNN iter:  
  for u in S:      No need to learn for-loops  
     $h_u^{(k)} = \sum_v \text{MLP}(h_v^{(k-1)}, h_u^{(k-1)})$ 
```

Bellman-Ford algorithm


```
for k = 1 ... |S| - 1:  
  for u in S:  
     $d[k][u] = \min_v d[k-1][v] + \text{cost}(v, u)$ 
```

Learns a simple reasoning step


 $h_u^{(k)}$
 $d[k][u]$
 \sum_v
 \min_v
 $\text{MLP}(h_v^{(k-1)}, h_u^{(k-1)})$
 $d[k-1][v] + \text{cost}(v, u)$

Neural exhaustive search


$\text{MLP}_2(\max_{\tau \subseteq S} \text{MLP}_1 \circ \text{LSTM}(X_1, \dots, X_{|\tau|} : X_1, \dots, X_{|\tau|} \in \tau)).$




(a) Maximum value difference.



(c) Monster trainer.




(b) Furthest pair.



(d) Subset sum. Random guessing yields 50%.

If alignment exists → step-by-step supervision

- Merely simulate the classical graph algorithm, generalizable
- No algorithm discovery



| Algorithm | Inputs | Supervision signals |
|----------------------|--|--|
| Breadth-first search | $x_i^{(t)}$: is i reachable from s in $\leq t$ hops? | $x_i^{(t+1)}$, $\tau^{(t)}$: has the algorithm terminated? |
| Bellman-Ford | $x_i^{(t)}$: shortest distance from s to i (using $\leq t$ hops) | $x_i^{(t+1)}$, $\tau^{(t)}$, $p_i^{(t)}$: predecessor of i in the shortest path tree (in $\leq t$ hops) |
| Prim's algorithm | $x_i^{(t)}$: is node i in the (partial) MST (built from s after t steps)? | $x_i^{(t+1)}$, $\tau^{(t)}$, $p_i^{(t)}$: predecessor of i in the partial MST |

Joint training is encouraged

Table 1: Accuracy of predicting reachability at different test-set sizes, trained on graphs of 20 nodes. GAT* correspond to the best GAT setup as per Section 3 (GAT-full using the full graph).

| Model | Reachability (mean step accuracy / last-step accuracy) | | |
|---------------------------------------|--|-------------------------------|------------------------|
| | 20 nodes | 50 nodes | 100 nodes |
| LSTM (Hochreiter & Schmidhuber, 1997) | 81.97% / 82.29% | 88.35% / 91.49% | 68.19% / 63.37% |
| GAT* (Veličković et al., 2018) | 93.28% / 99.86% | 93.97% / 100.0% | 92.34% / 99.97% |
| GAT-full* (Vaswani et al., 2017) | 78.40% / 77.86% | 85.76% / 91.83% | 88.98% / 91.51% |
| MPNN-mean (Gilmer et al., 2017) | 100.0% / 100.0% | 61.05% / 57.89% | 27.17% / 21.40% |
| MPNN-sum (Gilmer et al., 2017) | 99.66% / 100.0% | 94.25% / 100.0% | 94.72% / 98.63% |
| MPNN-max (Gilmer et al., 2017) | 100.0% / 100.0% | 100.0% / 100.0% | 99.92% / 99.80% |





Figure 3: The per-step algorithm execution performances in terms of reachability accuracy (**left**), distance mean-squared error (**middle**) and predecessor accuracy (**right**), tested on 100-node graphs after training on 20-node graphs. Please mind the scale of the MSE plot.

Processor as Transformer

- Back to input sequence (set), but stronger generalization
- Transformer with encoder mask ~ graph attention
- Use Transformer with:
 - Binary representation of numbers
 - Dynamic conditional masking



Training with execution trace

```
selection_sort(data):
    sorted_list = []
    while (len(data)) > 0:
        min_index, min_element = find_min(data)

        data.delete(min_index)
        sorted_list.append(min_element)
    return sorted_list

find_min(data):
    min_element = -1
    min_index = -1
    for index, element in enumerate(data):
        if (element < min_element):
            min_element = element
            min_index = index

    return [min_index, min_element]
```

```
merge_sort(data, start, end):
    if (start < end):
        mid = (start + end) / 2

        merge_sort(data, start, mid)
        merge_sort(data, mid+1, end)

    return merge(data, start, mid, end)
```

```
shortest_path(graph, source_node, shortest_path):
    dists = []
    nodes = []
    anchor_node = source_node
    node_list = graph.get_nodes()

    while node_list:
        possible_paths = sum(graph.adj(anchor_node),
                             shortest_path(anchor_node))
        shortest_path = min(possible_paths, shortest_path)

        anchor_node, min_dist = min(shortest_path(node_list))

        node_list.delete(anchor_node)
        nodes.append(anchor_node)
        dists.append(min_dist)

    return dists, nodes
```

```
minimum_spanning_tree(graph, source_node, node_val):
    mst_nodes = []
    mst_weights = []
    anchor_node = source_node
    res_nodes = graph.get_nodes()


    while node_list:
        adj_list = graph.adj(anchor_node)

        node_val(res_nodes) = min(node_val(res_nodes),
                                 adj_list(res_nodes))

        anchor_node, min_weight = min(node_val(res_nodes))

        mst_nodes.append(anchor_node)
        mst_weights.append(min_weight)
        res_nodes.delete(anchor_node)

    return mst_nodes, mst_weights
```



The results show strong generalization




Figure 5: 3D PCA visualization of learned bitwise embeddings for different numeric tasks. The embeddings exhibit regular, task-dependent structure, even when most numbers have not been seen in training (c).




Figure 3: Sorting performance of transformers trained on sequences of up to length 8.




Figure 4: Visualizing decoder attention weights. Attention is over each row. Transformer attention saturates as the output sequence length increases, while NEE maintains sharp attention.

| Sizes | 25 | 50 | 75 | 100 |
|-----------------------|--------|--------|--------|---------|
| Accuracy | | | | |
| Selection sort | 100.00 | 100.00 | 100.00 | 100.00 |
| Merge sort | 100.00 | 100.00 | 100.00 | 100.00 |
| Shortest path | 100.00 | 100.00 | 100.00 | 100.00* |
| Minimum spanning tree | 100.00 | 100.00 | 100.00 | 100.00 |

QA